

Sara Papi

✉ sara95papi@gmail.com ✉ @sarapapi ✉ sara-papi
🌐 <https://sarapapi.github.io/> ✉ Google Scholar (h-index 12)



Experience

- 2024 – now █ **AI Researcher** in Speech Processing and Multimodal LLMs (*Fondazione Bruno Kessler, Italy*) – WP Technical Lead in the MEETWEEN Horizon European Project and Task Lead in the DVPS Horizon European Project
- 2023 – 2023 █ **AI Research Intern** in Streaming Speech Recognition and Translation (*Microsoft, USA*)
– Supervised by Jinyu Li
- 2019 – 2019 █ **AI Research Intern** in Automatic Speech Scoring (*Fondazione Bruno Kessler, Italy*)

Education

- 2020 – 2024 █ **Ph.D. cum Laude** in Information Engineering and Computer Science (Fondazione Bruno Kessler & University of Trento, Italy)
- 2017 – 2020 █ **M.Sc. cum Laude** in Computer and Automation Engineering (University of Siena, Italy)
- 2014 – 2017 █ **B.Sc.** in Information Engineering (University of Siena, Italy)

Awards and Achievements

- 2025 █ **Best PhD Graduate 2024 in Information and Communication Technology** by University of Trento
- █ **PhD Excellence Award** by Fondazione Bruno Kessler
- █ **Highly Commended 2024 Anthony C Clarke Award** for the EAMT Best Thesis by EAMT
- 2024 █ **Special Mention for the Best Italian PhD Thesis** by AIxIA 2024.
- █ **Outstanding Paper and SAC Award** by ACL 2024
- █ **Social Impact Paper Award** by EMNLP 2024

Invited Talks/Keynotes

- 2026 █ **Cohere Labs Open Science Community.** Talk Title: “*Towards Crosslingual Evaluation of General-Purpose Instruction-Following Models Across Text, Speech, and Vision*”.
- 2025 █ **MT Marathon.** Talk Title: “*Speech Translation in the LLM Era: Between Old and New Challenges*”.
- █ **SlatorCon Remote.** Talk Title: “*How “Real” is your Real-Time Simultaneous Speech-to-Text System?*”.
- 2024 █ **IWSLT 2024.** Panel Title: “*Emerging trends in Speech Translation*”.

Research Publications

Journal Articles

- 1 █ **S. Papi**, P. Polák, D. Macháček, and O. Bojar, “How “real” is your real-time simultaneous speech-to-text translation system?” *Transactions of the Association for Computational Linguistics*, vol. 13, pp. 281–313, Apr. 2025, ISSN: 2307-387X. DOI: 10.1162/tacl_a_00740 eprint: https://direct.mit.edu/tacl/article-pdf/doi/10.1162/tacl_a_00740/2513610/tacl_a_00740.pdf.

- 2 S. Papi, M. Gaido, A. Karakanta, M. Cettolo, M. Negri, and M. Turchi, "Direct speech translation for automatic subtitling," *Transactions of the Association for Computational Linguistics*, vol. 11, pp. 1355–1376, Nov. 2023, ISSN: 2307-387X. DOI: 10.1162/tacl_a_00607

Conference Proceedings

- 1 I. Abdulmumin et al., "Findings of the IWSLT 2025 evaluation campaign," in *Proceedings of the 22nd International Conference on Spoken Language Translation (IWSLT 2025)*, Jul. 2025, ISBN: 979-8-89176-272-5. DOI: 10.18653/v1/2025.iwslt-1.44
- 2 M. Gaido*, S. Papi*, L. Bentivogli, A. Brutti, M. Cettolo, and other, "The warmup dilemma: How learning rate strategies impact speech-to-text model convergence," in *Proceedings of the 22nd International Conference on Spoken Language Translation (IWSLT 2025)*, Vienna, Austria (in-person and online), Jul. 2025, ISBN: 979-8-89176-272-5. DOI: 10.18653/v1/2025.iwslt-1.4
- 3 T. K. Lam*, M. Gaido*, S. Papi, L. Bentivogli, and B. Haddow, "Prepending or cross-attention for speech-to-text? an empirical comparison," in *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, Apr. 2025, ISBN: 979-8-89176-189-6. URL: <https://aclanthology.org/2025.naacl-long.153/>
- 4 N. Rao Koluguri et al., "Granary: Speech Recognition and Translation Dataset in 25 European Languages," in *Interspeech 2025*, 2025. DOI: 10.21437/Interspeech.2025-540
- 5 F. Verdini et al., "How to Connect Speech Foundation Models and Large Language Models? What Matters and What Does Not," in *Interspeech 2025*, 2025. DOI: 10.21437/Interspeech.2025-2245
- 6 M. Züfle, S. Papi, B. Savoldi, M. Gaido, L. Bentivogli, and J. Niehues, "NUTSHELL: A dataset for abstract generation from scientific talks," in *Proceedings of the 22nd International Conference on Spoken Language Translation (IWSLT 2025)*, Jul. 2025, ISBN: 979-8-89176-272-5. DOI: 10.18653/v1/2025.iwslt-1.2
- 7 I. S. Ahmad et al., "FINDINGS OF THE IWSLT 2024 EVALUATION CAMPAIGN," in *Proceedings of the 21st International Conference on Spoken Language Translation (IWSLT 2024)*, Aug. 2024. DOI: 10.18653/v1/2024.iwslt-1.1
- 8 M. Cettolo, A. Piergentili, S. Papi, M. Gaido, M. Negri, L. Bentivogli, et al., "Magnet-machines generating translations: A calamita challenge," in *Proceedings of the Tenth Italian Conference on Computational Linguistics (CLiC-it 2024)*, 2024.
- 9 M. Gaido, S. Papi, M. Negri, and L. Bentivogli, "Speech translation with speech foundation models and large language models: What is there and what is missing?" In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Aug. 2024. DOI: 10.18653/v1/2024.acl-long.789
- 10 M. Gaido, S. Papi, M. Negri, M. Cettolo, and L. Bentivogli, "SBAAM! eliminating transcript dependency in automatic subtitling," in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Aug. 2024. DOI: 10.18653/v1/2024.acl-long.201
- 11 M. Gaido et al., "Automatic subtitling and subtitle compression: FBK at the IWSLT 2024 subtitling track," in *Proceedings of the 21st International Conference on Spoken Language Translation (IWSLT 2024)*, Aug. 2024. DOI: 10.18653/v1/2024.iwslt-1.13
- 12 M. Gaido*, S. Papi*, L. Bentivogli, A. Brutti, M. Cettolo, and other, "MOSEL: 950,000 hours of speech data for open-source speech foundation model training on EU languages," in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, Nov. 2024. DOI: 10.18653/v1/2024.emnlp-main.771

- 13 M. Gaido*, **S. Papi***, M. Negri, and L. Bentivogli, "How do hyenas deal with human speech? speech recognition and translation with ConfHyena," in *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, May 2024. ↗ URL: <https://aclanthology.org/2024.lrec-main.717/>
- 14 **S. Papi**, M. Gaido, M. Negri, and L. Bentivogli, "SimulSeamless: FBK at IWSLT 2024 simultaneous speech translation," in *Proceedings of the 21st International Conference on Spoken Language Translation (IWSLT 2024)*, Aug. 2024. ↗ DOI: 10.18653/v1/2024.iwslt-1.11
- 15 **S. Papi**, M. Gaido, M. Negri, and L. Bentivogli, "StreamAtt: Direct streaming speech-to-text translation with attention-based audio history selection," in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Aug. 2024. ↗ DOI: 10.18653/v1/2024.acl-long.202
- 16 **S. Papi** et al., "Leveraging timestamp information for serialized joint streaming recognition and translation," in *ICASSP 2024 - 2024 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2024. ↗ DOI: 10.1109/ICASSP48485.2024.10447565
- 17 **S. Papi***, M. Gaido*, A. Pilzer, and M. Negri, "When good and reproducible results are a giant with feet of clay: The importance of software quality in NLP," in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Aug. 2024. ↗ DOI: 10.18653/v1/2024.acl-long.200
- 18 B. Savoldi, **S. Papi**, M. Negri, A. Guerberof-Arenas, and L. Bentivogli, "What the harm? quantifying the tangible impact of gender bias in machine translation with a human-centered study," in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, Nov. 2024. ↗ DOI: 10.18653/v1/2024.emnlp-main.1002
- 19 D. Fucci, M. Gaido, **S. Papi**, M. Cettolo, M. Negri, and L. Bentivogli, "Integrating language models into direct speech translation: An inference-time solution to control gender inflection," in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, Dec. 2023. ↗ DOI: 10.18653/v1/2023.emnlp-main.705
- 20 M. Gaido, **S. Papi**, M. Negri, and M. Turchi, "Joint speech translation and named entity recognition," in *Interspeech 2023*, 2023. ↗ DOI: 10.21437/Interspeech.2023-1767
- 21 **S. Papi**, M. Gaido, and M. Negri, "Direct models for simultaneous translation and automatic subtitling: FBK@IWSLT2023," in *Proceedings of the 20th International Conference on Spoken Language Translation (IWSLT 2023)*, Jul. 2023. ↗ DOI: 10.18653/v1/2023.iwslt-1.11
- 22 **S. Papi**, M. Negri, and M. Turchi, "Attention as a guide for simultaneous speech translation," in *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Jul. 2023. ↗ DOI: 10.18653/v1/2023.acl-long.745
- 23 **S. Papi**, M. Turchi, and M. Negri, "Alignatt: Using attention-based audio-translation alignments as a guide for simultaneous speech translation," in *Interspeech 2023*, 2023. ↗ DOI: 10.21437/Interspeech.2023-170
- 24 **S. Papi**, P. Wang, J. Chen, J. Xue, J. Li, and Y. Gaur, "Token-level serialized output training for joint streaming asr and st leveraging textual alignments," in *2023 IEEE Automatic Speech Recognition and Understanding Workshop (ASRU)*, 2023. ↗ DOI: 10.1109/ASRU57964.2023.10389715
- 25 M. Gaido*, **S. Papi***, D. Fucci, G. Fiameni, M. Negri, and M. Turchi, "Efficient yet competitive speech translation: FBK@IWSLT2022," in *Proceedings of the 19th International Conference on Spoken Language Translation (IWSLT 2022)*, May 2022. ↗ DOI: 10.18653/v1/2022.iwslt-1.13
- 26 **S. Papi**, M. Gaido, M. Negri, and M. Turchi, "Does simultaneous speech translation need simultaneous models?" In *Findings of the Association for Computational Linguistics: EMNLP 2022*, Dec. 2022. ↗ DOI: 10.18653/v1/2022.findings-emnlp.11

- 27 **S. Papi**, M. Gaido, M. Negri, and M. Turchi, “Over-generation cannot be rewarded: Length-adaptive average lagging for simultaneous speech translation,” in *Proceedings of the Third Workshop on Automatic Simultaneous Translation*, Jul. 2022.  DOI: 10.18653/v1/2022.autosimtrans-1.2
- 28 **S. Papi**, A. Karakanta, M. Negri, and M. Turchi, “Dodging the data bottleneck: Automatic subtitling with automatically segmented ST corpora,” in *Proceedings of the 2nd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics and the 12th International Joint Conference on Natural Language Processing (Volume 2: Short Papers)*, Nov. 2022.  DOI: 10.18653/v1/2022.aacl-short.59
- 29 A. Karakanta*, **S. Papi***, M. Negri, and M. Turchi, “Simultaneous speech translation for live subtitling: From delay to display,” in *Proceedings of the 1st Workshop on Automatic Spoken Language Translation in Real-World Settings (ASLTRW)*, Aug. 2021.  URL: <https://aclanthology.org/2021.mtsummit-asltrw.4/>
- 30 **S. Papi**, M. Gaido, M. Negri, and M. Turchi, “Dealing with training and test segmentation mismatch: FBK@IWSLT2021,” in *Proceedings of the 18th International Conference on Spoken Language Translation (IWSLT 2021)*, Aug. 2021.  DOI: 10.18653/v1/2021.iwslt-1.8
- 31 **S. Papi**, M. Negri, and M. Turchi, “Visualization: The missing factor in simultaneous speech translation,” in *Proceedings of the Seventh Italian Conference on Computational Linguistics (CLiC-it 2021)*, 2021.
- 32 **S. Papi***, M. Gaido*, M. Negri, and M. Turchi, “Speechformer: Reducing information loss in direct speech translation,” in *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, Nov. 2021.  DOI: 10.18653/v1/2021.emnlp-main.127
- 33 **S. Papi**, E. Trentin, R. Gretter, M. Matassoni, and D. Falavigna, “Mixtures of deep neural experts for automated speech scoring,” in *Interspeech 2020*, 2020.  DOI: 10.21437/Interspeech.2020-1055

Under Review

- 1 M. Cettolo, M. Gaido, M. Negri, **S. Papi**, and L. Bentivogli, *How to evaluate speech translation with source-aware neural mt metrics*, 2025. arXiv: 2511.03295 [cs.CL].  URL: <https://arxiv.org/abs/2511.03295>
- 2 M. Gaido, **S. Papi**, M. Cettolo, M. Negri, and L. Bentivogli, *Simulstream: Open-source toolkit for evaluation and demonstration of streaming speech-to-text translation systems*, 2025. arXiv: 2512.17648 [cs.CL].  URL: <https://arxiv.org/abs/2512.17648>
- 3 M. Nissim et al., *Challenging the abilities of large language models in italian: A community initiative*, 2025. arXiv: 2512.04759 [cs.CL].  URL: <https://arxiv.org/abs/2512.04759>
- 4 **S. Papi**, D. Fucci, M. Gaido, M. Negri, and L. Bentivogli, *Cross-attention is half explanation in speech-to-text models*, 2025. arXiv: 2509.18010 [cs.CL].  URL: <https://arxiv.org/abs/2509.18010>
- 5 **S. Papi** et al., *Hearing to translate: The effectiveness of speech modality integration into llms*, 2025. arXiv: 2512.16378 [cs.CL].  URL: <https://arxiv.org/abs/2512.16378>
- 6 **S. Papi** et al., *MCIF: Multimodal crosslingual instruction-following benchmark from scientific talks*, 2025. arXiv: 2507.19634 [cs.CL].  URL: <https://arxiv.org/abs/2507.19634>
- 7 P. Polák, **S. Papi**, L. Bentivogli, and O. Bojar, *Better late than never: Evaluation of latency metrics for simultaneous speech-to-text translation*, 2025. arXiv: 2509.17349 [cs.CL].  URL: <https://arxiv.org/abs/2509.17349>

Community Service

Organizer/Committee Member

- 2026  **IWSLT 2026 - Evaluation Campaigns**, Instruction Following (Chair) and Simultaneous Translation (Organizer).
- 2025  **IWSLT 2025 - Evaluation Campaigns**, Instruction Following (Chair) and Simultaneous Translation (Organizer).
- 2024  **EACL 2024 Student Research Workshop** (Student Program Chair). Conference.
-  **IWSLT 2024 - Evaluation Campaigns**, Subtitling (Organizer). Conference.

Area Chair/Reviewer

- 2026  **LANS-TTS - IEEE Access** (Journal)
- 2025  **ACL, EMNLP** (Conferences) – Area Chair
  **NeurIPS, CLIC-it, Interspeech, IWSLT, ICASSP 2025 (SPADE Workshop)** (Conferences) – Reviewer
  **LANS-TTS - IEEE Access** (Journal)
- 2024  **ACL, EMNLP, ICASSP, IWSLT, Interspeech** (Conferences) – Reviewer
- 2023  **HITLAML** (Conference) – Reviewer
- 2022  **AMTA 2022** (Conference) – Reviewer